

Foundations of Data Science Fall 2022 - Homework 2 (30 points)

Student Name: Yamini Lakshmi Narasimhan

Student Net Id: yl9822

Part 1: Preparing a Training Set and Training a Decision Tree (10 Points)

This is a hands-on task where we build a predictive model using Decision Trees discussed in class. For this part, we will be using the data in `cell2cell_data.csv` (you can find this on NYU Brightspace).

These historical data consist of 39,859 customers: 19,901 customers that churned (i.e., left the company) and 19,958 that did not churn (see the `"churndep"` variable). Here are the data set's 11 possible predictor variables for churning behavior:

Pos.	Var. Name	Var. Description
-----	-----	-----
1	revenue	Mean monthly revenue in dollars
2	outcalls	Mean number of outbound voice calls
3	incalls	Mean number of inbound voice calls
4	months	Months in Service
5	eqpdays	Number of days the customer has had his/her current equipment
6	webcap	Handset is web capable
7	marryyes	Married (1=Yes; 0=No)
8	travel	Has traveled to non-US country (1=Yes; 0=No)
9	pcown	Owns a personal computer (1=Yes; 0=No)
10	creditcd	Possesses a credit card (1=Yes; 0=No)
11	retcalls	Number of calls previously made to retention team

The 12th column, the dependent variable `"churndep"`, equals 1 if the customer churned, and 0 otherwise.

```
In [202... import warnings
from pprint import pprint
warnings.filterwarnings('ignore')
warnings.filterwarnings(action='once')
```

1. Load the data and prepare it for modeling. Note that the features are already processed for you, so the only thing needed here is split the data into training and testing. Use pandas to create two data frames: `train_df` and `test_df`, where `train_df` has 80% of the data chosen uniformly at random without replacement (`test_df` should have the other 20%). Also, make sure to write your own code to do the splits. You may use any `random()` function numpy but do not use the data splitting functions from Sklearn.

(2 Points)

```
In [165... # Place your code here
import pandas as pd
import numpy as np

df = pd.read_csv("./cell2cell_data.csv")

print("Total data", len(df))
percent = np.random.rand(len(df)) <= 0.8
train_df = df[percent]
test_df = df[~percent]

print("Train df", len(train_df))
print("Test df", len(test_df))

train_X = train_df.loc[:, train_df.columns != "churndep"]
train_Y = train_df["churndep"]

test_X = test_df.loc[:, test_df.columns != "churndep"]
test_Y = test_df["churndep"]
```

Total data 39833

Train df 31910

Test df 7923

```
In [166... print(train_df["churndep"].value_counts())
print(test_df["churndep"].value_counts())
```

```

0    15985
1    15925
Name: churndep, dtype: int64
0     3965
1     3958
Name: churndep, dtype: int64

```

In [167... `df.head()`

```

Out[167]:
   revenue  outcalls  incalls  months  eqpdays  webcap  marryyes  travel  pdown  creditc
0    48.82    10.00     3.0     26     780      0         0         0         0         0
1    83.53    20.00     1.0     31     745      1         0         0         0         0
2    29.99     0.00     0.0     52    1441      0         0         0         1         0
3    51.42     0.00     0.0     36      59      1         0         0         0         0
4    37.75     2.67     0.0     25     572      0         0         0         1         0

```

2. Now build and train a decision tree classifier using `DecisionTreeClassifier()` ([manual page](#)) on `train_df` to predict the `"churndep"` target variable. Make sure to use `criterion='entropy'` when instantiating an instance of `DecisionTreeClassifier()`. For all other settings you should use all of the default options.

(1 Point)

```

In [168... # Place your code here
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier(criterion='entropy')
dtc.fit(train_X, train_Y)

```

```

Out[168]:
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')

```

3. Using the resulting model from 1.2, show a bar plot of feature names and their feature importance (hint: check the attributes of the `DecisionTreeClassifier()` object directly in IPython or check the manual!).

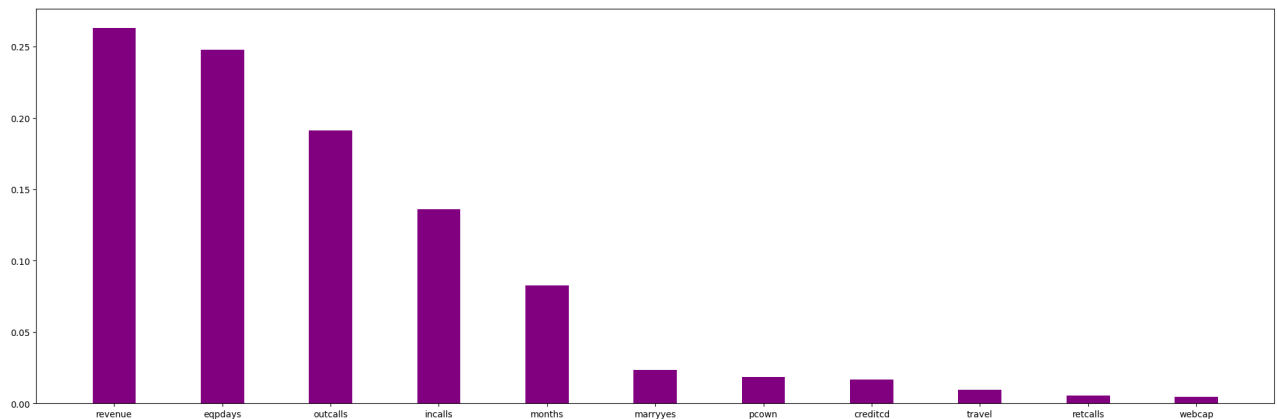
(3 Points)

```
In [169]: # Place your code here
import matplotlib.pyplot as plt
feature_importance = {}
for importance, name in sorted(zip(dtc.feature_importances_, train_X), reverse=True):
    feature_importance[name] = importance

plt.figure(figsize=(25,8))
print(feature_importance)
plt.bar(feature_importance.keys(), feature_importance.values(), color='purple')

{'revenue': 0.2633111890880889, 'eqpdays': 0.2477853547757061, 'outcalls': 0.19113101270417757, 'incalls': 0.13606175408566842, 'months': 0.08289865802309869, 'marryyes': 0.02353186425352615, 'pcown': 0.01855427418985122, 'creditcd': 0.01671943163250753, 'travel': 0.009737491800810492, 'retcalls': 0.005654106620520814, 'webcap': 0.004614862826043981}
```

Out[169]: <BarContainer object of 11 artists>



4. Is the relationship between the top 3 most important features (as measured here) negative or positive? If your marketing director asked you to explain the top 3 drivers of churn, how would you interpret the relationship between these 3 features and the churn outcome? What "real-life" connection can you draw between each variable and churn?

(2 Points)

```
In [112]: df_corr = df[["revenue", "eqpdays", "outcalls", "churndep"]]
df_corr.corr()
```

Out[112]:

	revenue	eqpdays	outcalls	churndep
revenue	1.000000	-0.222074	0.500709	-0.013370
eqpdays	-0.222074	1.000000	-0.244112	0.112821
outcalls	0.500709	-0.244112	1.000000	-0.037071
churndep	-0.013370	0.112821	-0.037071	1.000000

Top 3 drivers of churn are revenue, outcalls and eqpdays of the company revenue <-> churndep -0.013370 negative outcalls<-> churndep -0.037071 negative eqpdays <-> churndep +0.112821 positive

Revenue is negatively correlated that means higher the revenue (Mean monthly revenue in dollars) lower the churndep, similarly outcalls and churndep are also negatively correlated, eqpdays and churndep are positively correlated that is higher eqpdays then higher churndep

Revenue: Revenue is negatively correlated that means higher the revenue lower the churndep It makes sense that if the revenue is high for the company then churndep is low as the clients will trust the company lot more since profits are flowing in Another relation could be that if revenue is flowing in it means clients are flushing the money in so high revenue means low number of clients are leaving and lot of clients are currently paying

Outcalls: Outcalls being negatively correlated also makes sense as it means Mean number of outbound voice calls with the client is lot so client feels a lot more reliable of the company so churndep is lower.

Eqpdays: Number of days the customer has had his/her current equipment : positively correlated which means higher the no of days the customer has their equipment, higher the churndep. May be the equipment becomes faulty after few days and there are no proper customer service because of which the customer tends to leave.

5. Using the classifier built in 1.2, try predicting "churndep" on both the train_df and test_df data sets. What is the accuracy on each?

(2 Points)

In [170...

```
# Place your code here
# Place your code here
from sklearn.metrics import accuracy_score
predtest_Y = dtc.predict(test_X)
print("Test Accuracy", accuracy_score(test_Y, predtest_Y))

predtrain_Y = dtc.predict(train_X)
print("Train Accuracy", accuracy_score(train_Y, predtrain_Y))
```

Test Accuracy 0.5367916193361101
Train Accuracy 0.9998433093074272

Part 2 - Finding a Good Decision Tree (Total 10 Points)

The default options for your decision tree may not be optimal. We need to analyze whether tuning the parameters can improve the accuracy of the classifier. For the following options `min_samples_split` and `min_samples_leaf`:

1. Generate a list of 10 values of each for the parameters `min_samples_split` and `min_samples_leaf`.

(1 Point)

`min_samples_leaf`: int, float, optional (default=1)
The minimum number of samples required to be at a leaf node. A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches. This may have the effect of smoothing the model, especially in regression. `min_samples_leaf` is also used to control over-fitting by defining that each leaf has more than one element. Thus ensuring that the tree cannot overfit the training dataset by creating a bunch of small branches exclusively for one sample each. In reality, what this is actually doing is simply just telling the tree that each leaf doesn't have to have an impurity of 0

`min_samples_split`: int or float, default=2
The minimum number of samples required to split an internal node:
If int, then consider `min_samples_split` as the minimum number.
If float, then `min_samples_split` is a fraction and `ceil(min_samples_split * n_samples)` are the minimum number of samples for each split.

2. Explain in words your reasoning for choosing the above ranges.

`min_samples_leaf`

Chose a range between 100–1750 with a difference of 250 because since we have a dataset of 22k and only 2 classes to predict with only 5 features that contribute to the predictions – revenue, eqpdays, outcalls, incalls and months – the child nodes definitely can have a minimum of range 100 and can go higher as the dataset will be majorly split by 5 features split – this will also prevent overfitting and generalising the model

`min_samples_split`

Chose the same range as we don't want the dataset to overfit on the training sample and having the less than 100 would not only overfit but will also create a complex tree with multiple nodes and for 5 contributing features a shallow tree would be sufficient with multiple samples grouped together in a node

Place your response here

3. For each combination of values in 3.1 (there should be 100), build a new classifier and check the classifier's accuracy on the test data. Plot the test set accuracy for these options. Use the values of `min_samples_split` as the x-axis and generate a new series (line) for each of `min_samples_leaf`.

(5 Points)

```

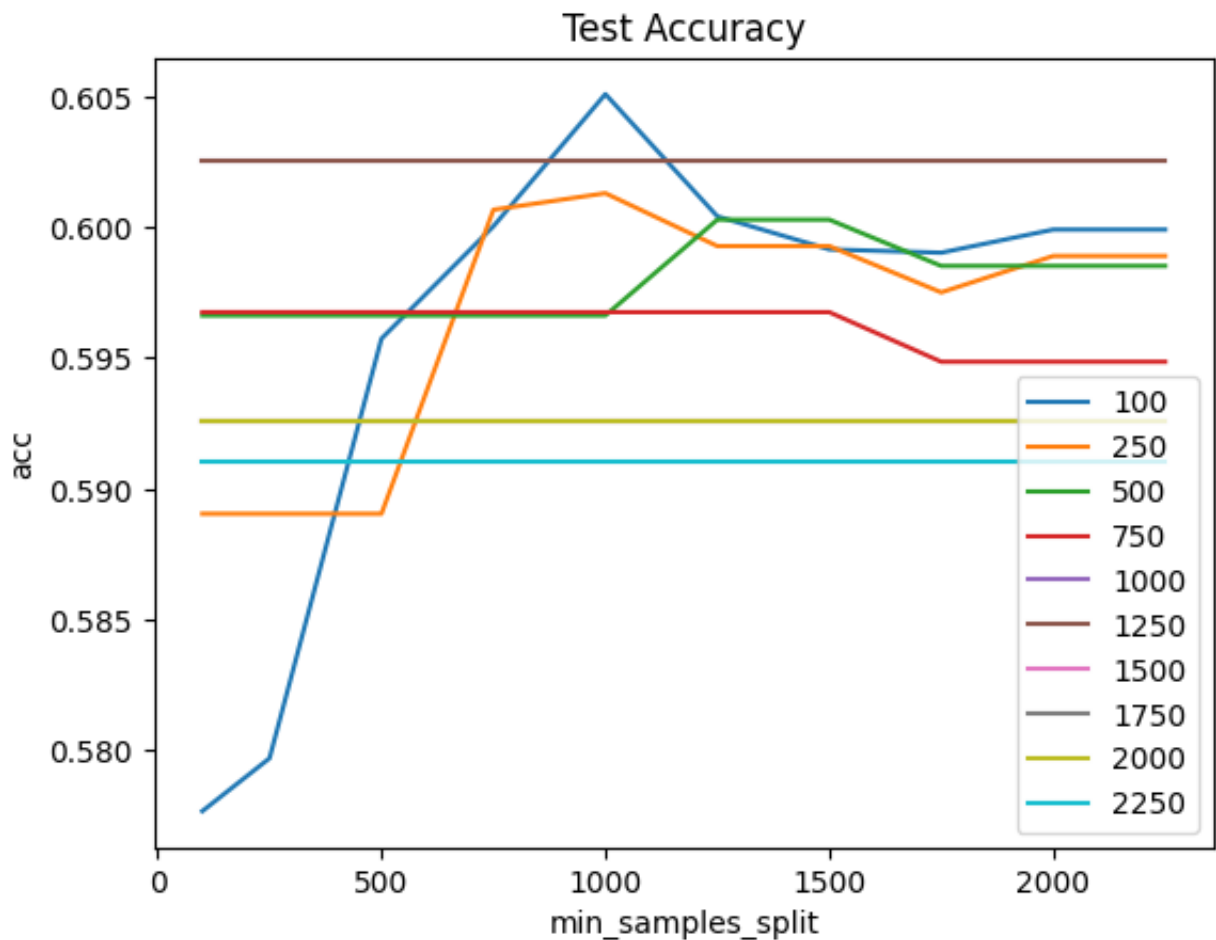
In [264... # Place your code here
min_samples_split = [100, 250, 500, 750, 1000, 1250, 1500, 1750, 2000, 2250]
min_samples_leaf = [100, 250, 500, 750, 1000, 1250, 1500, 1750, 2000, 2250]
j=0

for i in range(len(min_samples_leaf)):
    acc_lis = []
    for j in range(len(min_samples_split)):
        dtc = DecisionTreeClassifier(random_state=0, min_samples_split=min_s
        pretest_Y = dtc.predict(test_X)
        acc = accuracy_score(test_Y, pretest_Y)
        acc_lis.append(acc)

    line = plt.plot(min_samples_split, acc_lis, label = str(min_samples_leaf

plt.xlabel("min_samples_split")
plt.ylabel("acc")
plt.legend()
plt.title("Test Accuracy")
plt.show()

```



4. Which configuration returns the best accuracy? What is this accuracy? (Note, if you don't see much variation in the test set accuracy across values of `min_samples_split` or `min_samples_leaf`, try redoing the above steps with a different range of values).

(1 Point)

Best configuration is when `min_samples_split` is 1000 and `min_samples_leaf` is 100 and accuracy is around 60.5%

Place your response here

5. If you were working for a marketing department, how would you use your churn production model in a real business environment? Explain why churn prediction might be good for the business and how one might improve churn by using this model.

(2 Points)

I would use this model to periodically check which customer/client might have a tendency to get churned and take precautionary measure beforehand. For example if the model predicts a client/customer is going to leave and then we can always increase outbound and inbound calls as that might make the customer feel they're appreciated and stay with the company for longer or bring in more revenue by transparently showing them how we can improve

I can change my focus to companies that have a greater tendency to drop working with us there by giving them a better service and better deal for them to stay with us longer

Place your response here

Part 3: Model selection with cross-validation (5 points)

In this part, we will focus on cross-validation to find a good value for parameter `max_depth`.

1. Write a cross-validation function that does the following:

- Takes as inputs a dataset, a label name, # of splits/folds (`k`), and a sequence of values for the maximum depth of the tree (`max_depth`).
- Shuffles the data.
- Splits the data into `k` folds according to the cross-validation logic
- Performs two loops
 - Outer Loop: `for each f in range(k) :`
 - Inner Loop: `for each value in max_depth_sequence :`
 - Trains a Decision Tree on the training split with the `max_depth=value` (USE `criterion='entropy'` BUT DO NOT ALTER THE OTHER PARAMETERS)
 - Computes `accuracy_value_f` on test split
 - Stores `accuracy_value_f` in a dictionary of values
- Returns a dictionary, where each key-value pair is: `value: [accuracy_value_1,...,accuracy_value_k]`

(2 Points)

In [210... **from** random **import** randrange

```
def xValDecisionTree(dataset, label, k, max_depth_sequence):
    dataset = dataset.sample(frac=1).reset_index(drop=True)
    df_split = np.array_split(dataset, k)
    train_split = pd.DataFrame()
    acc_dict = {}
    for i in range(0,k):
        test_split = pd.DataFrame(df_split[i])
        for j in range(k):
            if i!= j:
                train_split = pd.concat([pd.DataFrame(df_split[j]), train_split])
        acc_lis = []
        depth_lis = []

        train_X = train_split.loc[:, train_split.columns != label]
        train_Y = train_split[label]

        test_X = test_split.loc[:, test_split.columns != label]
        test_Y = test_split[label]

        for value in max_depth_sequence:
            dtc = DecisionTreeClassifier(random_state=0, max_depth=value, c
            pretest_Y = dtc.predict(test_X)
            acc = accuracy_score(test_Y, pretest_Y)
            acc_lis.append(acc)
            depth_lis.append(value)
            if value not in acc_dict.keys():
                acc_dict[value] = []
            acc_dict[value].append(acc)
        plt.show()

    return acc_dict
```

2. Using the function written above, do the following:

- Generate a sequence `max_depth_sequence = [None, 2, 4, 8, 16, 32, 128, 256, 512]` (Note that None is the default value for this parameter).
- 2. Call `accs = xValDecisionTree(dataset, 'churndep', 10, max_depth_sequence)`
- 3. For each value in `accs.keys()`, calculate `mean(accs[value])`. What value is associated with the highest accuracy mean?
- 4. For each value in `accs.keys()`, calculate the ranges `mean(accs[value]) +/- std(accs[value])`. Do the ranges associated with the value that has the highest `mean(accs[value])` overlap with ranges for other values? What may this suggest and what are the limitations of a standard deviation based analysis in this scenario?
- 5. Which depth value would you pick, if any, and why?

(3 Points)

```
In [257... # Place your code here
#2.
max_depth_sequence = [None, 2, 4, 8, 16, 32, 128, 256, 512]
accs_dict = xValDecisionTree(df, "churndep", 10, max_depth_sequence)
pprint(accs_dict)
```

```
{None: [0.5348895582329317,
        0.9997489959839357,
        1.0,
        0.9992467988953051,
        0.9997489329651017,
        0.9997489329651017,
        0.9997489329651017,
        0.9997489329651017,
        1.0,
        0.9997489329651017],
 2: [0.5820783132530121,
      0.5747991967871486,
      0.5860943775100401,
      0.56866683404469,
      0.5849861913130806,
      0.5824755209640974,
      0.5827265879989957,
      0.5895053979412503,
      0.5827265879989957,
      0.590760733115742],
 4: [0.5978915662650602,
      0.5956325301204819,
```

```
0.5938755020080321,  
0.5842329902083856,  
0.5985438111975897,  
0.5872457946271654,  
0.5992970123022847,  
0.6055736881747427,  
0.6008034145116746,  
0.5967863419533015],  
8: [0.5946285140562249,  
0.6056726907630522,  
0.6054216867469879,  
0.6038162189304545,  
0.6136078332914888,  
0.6010544815465729,  
0.6171227717800652,  
0.6148631684659804,  
0.6211398443384384,  
0.6048204870700477],  
16: [0.5662650602409639,  
0.6859939759036144,  
0.696285140562249,  
0.6921918152146623,  
0.6869194074817977,  
0.6768767260858649,  
0.6778809942254582,  
0.7102686417273412,  
0.7062515691689681,  
0.7072558373085613],  
32: [0.5431726907630522,  
0.9427710843373494,  
0.9492971887550201,  
0.9540547326136078,  
0.9450163193572684,  
0.950288727090133,  
0.9480291237760482,  
0.9422545819733869,  
0.9590760733115742,  
0.9477780567411499],  
128: [0.5348895582329317,  
0.9997489959839357,  
1.0,  
0.9992467988953051,  
0.9997489329651017,  
0.9997489329651017,  
0.9997489329651017,  
0.9997489329651017,  
1.0,  
0.9997489329651017],  
256: [0.5348895582329317,  
0.9997489959839357,
```

```

1.0,
0.9992467988953051,
0.9997489329651017,
0.9997489329651017,
0.9997489329651017,
0.9997489329651017,
1.0,
0.9997489329651017],
512: [0.5348895582329317,
0.9997489959839357,
1.0,
0.9992467988953051,
0.9997489329651017,
0.9997489329651017,
0.9997489329651017,
0.9997489329651017,
0.9997489329651017,
1.0,
0.9997489329651017]]}

```

```
In [214]: accs_df = pd.DataFrame.from_dict(accs_dict)
accs_df
```

```
Out[214]:
```

	NaN	2.0	4.0	8.0	16.0	32.0	128.0	256.0	
0	0.533133	0.577309	0.586847	0.581576	0.565261	0.535392	0.533133	0.533133	0.5
1	0.999749	0.590110	0.604418	0.609940	0.694528	0.948042	0.999749	0.999749	0.9
2	0.999749	0.581576	0.594378	0.601155	0.693273	0.943775	0.999749	0.999749	0.9
3	0.999749	0.582224	0.596033	0.608838	0.684911	0.934220	0.999749	0.999749	0.9
4	0.998745	0.581973	0.595280	0.607331	0.687170	0.934220	0.998745	0.998745	0.9
5	1.000000	0.576199	0.592769	0.607080	0.691690	0.949033	1.000000	1.000000	1.0
6	0.999749	0.590008	0.603314	0.620889	0.697464	0.950540	0.999749	0.999749	0.9
7	0.999749	0.576450	0.591765	0.607582	0.693698	0.939744	0.999749	0.999749	0.9
8	1.000000	0.586744	0.600803	0.618880	0.712779	0.947527	1.000000	1.000000	1.0
9	0.999498	0.580216	0.590761	0.603565	0.696711	0.945518	0.999498	0.999498	0.9

The maximum accuracy in the val set is 0.999498 when max_depth = None, 128, 256, 512

```
In [256]: #3
describe_df = accs_df.describe()
describe_df.loc['mean-std'] = describe_df.loc['mean'] - describe_df.loc['std']
describe_df.loc['mean+std'] = describe_df.loc['mean'] + describe_df.loc['std']
describe_df
```

Out [256]:

	NaN	2.0	4.0	8.0	16.0	32.0	128.0
count	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000
mean	0.953012	0.582281	0.595637	0.606684	0.681749	0.902801	0.953012
std	0.147531	0.005180	0.005665	0.010739	0.041607	0.129227	0.147531
min	0.533133	0.576199	0.586847	0.581576	0.565261	0.535392	0.533133
25%	0.999561	0.578036	0.592016	0.604444	0.688300	0.935601	0.999561
50%	0.999749	0.581775	0.594829	0.607457	0.693486	0.944647	0.999749
75%	0.999749	0.585614	0.599611	0.609664	0.696165	0.947913	0.999749
max	1.000000	0.590110	0.604418	0.620889	0.712779	0.950540	1.000000
mean-std	0.805481	0.577101	0.589972	0.595945	0.640142	0.773574	0.805481
mean+std	1.100543	0.587461	0.601302	0.617422	0.723356	1.032028	1.100543

4. Standard deviation here shows how the data in folds are distributed that is if the std of accuracy is low then it means that the data is distributed uniformly among the folds whereas if std is high then it shows that the data in the val fold is highly variable and is not similar to the training data folds

5. The range(mean - std to mean + std) is the same for depth = None, 256 and 512 its the same range that is 0.805481 - 1.100543

It means the data is overfitting to the train data so the val set accuracy is extremely high So, i'll probably take the next best accuracy that is with value 16.0 as max_depth_sequence

```
In [258... dtc = DecisionTreeClassifier(random_state=0, max_depth=None, criterion='entropy')
predtest_Y = dtc.predict(test_X)
acc = accuracy_score(test_Y, predtest_Y)
print("For max_depth = None Acc = ",acc)

dtc = DecisionTreeClassifier(random_state=0, max_depth=256, criterion='entropy')
predtest_Y = dtc.predict(test_X)
acc = accuracy_score(test_Y, predtest_Y)
print("For max_depth = 256 Acc = ",acc)

dtc = DecisionTreeClassifier(random_state=0, max_depth=512, criterion='entropy')
predtest_Y = dtc.predict(test_X)
acc = accuracy_score(test_Y, predtest_Y)
print("For max_depth = 512 Acc = ",acc)

For max_depth = None Acc = 0.5386848416004039
For max_depth = 256 Acc = 0.5386848416004039
For max_depth = 512 Acc = 0.5386848416004039
```

This shows how the test data is not performing that great on the best performing model in the train data as those iteration are over fitting, whereas the one with max_depth of 16 has generalised well so performs better than max_depth = None, 128, 256, 512

```
In [259... dtc = DecisionTreeClassifier(random_state=0, max_depth=16, criterion='entropy')
predtest_Y = dtc.predict(test_X)
acc = accuracy_score(test_Y, predtest_Y)
print("For max_depth = 16 Acc = ",acc)

For max_depth = 16 Acc = 0.566578316294333
```

In []:

Part 4: Boosting (5 Points)

Now, as we covered in class, ensemble methods are often used to improve performance.

1. Implement the boosting algorithm: XGBoost for the same `cell2cell_data.csv` task as above. You will have to select how to tune hyperparameters. Besides depth, which other hyperparameters do you optimize for? (2 points)

XGBoost hyperparameters

General Parameters

booster

gbtree default which is okay as it is used for tree based

xgboost

Booster Parameters

eta

It is the step size shrinkage used in update to prevent overfitting.

Range : $[0,1]$ Typical final values : 0.01-0.2.\n

Default : 0

gamma

A node is split only when the resulting split gives a positive reduction in the loss function. The larger gamma is, the more conservative the algorithm will be.

Range: $[0,\infty]$

Default : 0

max_depth – used 3 because mostly there are 3 major features that is eqpdays, revenue, outcalls

It is used to control over-fitting as higher depth will allow model to learn relations very specific to a particular sample.

range: $[0,\infty]$ (0 is only accepted in lossguided growing policy when tree_method is set as hist.

default =6

min_child_weight – for overfitting we have modified tree depth etc

It defines the minimum sum of weights of all observations required in a child.

range: $[0,\infty]$

It is used to control over-fitting.

Higher values prevent a model from learning relations which might be highly specific to the particular sample selected for a tree.

default=1

max_delta_step – churndep is not biased(0: 15991 1: 15951) so not using this

In maximum delta step we allow each tree's weight estimation to be.

Usually this parameter is not needed, but it might help in logistic regression when class is extremely imbalanced.

Set it to value of 1-10 might help control the update.

range: $[0,\infty]$

default=0

subsample – Want to give this default because this is too much rules on the tree

It denotes the fraction of observations to be randomly samples for each tree.

Subsample ratio of the training instances.

Setting it to 0.5 means that XGBoost would randomly sample half of the training data prior to growing trees. – This will prevent overfitting.

Typical values: 0.5–1

range: (0,1]

lambda – 1 is good enough

L2 regularization term on weights (analogous to Ridge regression).

This is used to handle the regularization part of XGBoost.

Increasing this value will make model more conservative.

default=1

alpha – No need this because less features and we don't need feature sleection here

L1 regularization term on weights (analogous to Lasso regression).

It can be used in case of very high dimensionality so that the algorithm runs faster when implemented.

Increasing this value will make model more conservative.

default=0

tree_method – auto takes care of which one to choose based on dataset size so no need of change

Choices: auto, exact, approx, hist, gpu_hist

auto: Use heuristic to choose the fastest method.

For small to medium dataset, exact greedy (exact) will be used.

For very large dataset, approximate algorithm (approx) will be chosen.

Because old behavior is always use exact greedy in single machine, user will get a message when approximate algorithm is chosen to notify this choice.

exact: Exact greedy algorithm.

approx: Approximate greedy algorithm using quantile sketch and gradient histogram.

hist: Fast histogram optimized approximate greedy algorithm.

It uses some performance improvements such as bins caching.
 gpu_hist: GPU implementation of hist algorithm.
 default= auto

scale_pos_weight – No need as classes are balanced for us
 It controls the balance of positive and negative weights,
 It is useful for imbalanced classes.
 A value greater than 0 should be used in case of high class
 imbalance as it helps in faster convergence.
 A typical value to consider: $\text{sum}(\text{negative instances}) / \text{sum}(\text{positive instances})$.
 default=1

max_leaves – too much interference to the tree so no need
 Maximum number of nodes to be added.
 Only relevant when grow_policy=lossguide is set.

Learning Task Parameters

objective – binary:logistic because this is binary
 classification
 reg:logistic : logistic regression
 binary:logistic : logistic regression for binary
 classification, output probability
 binary:logitraw: logistic regression for binary
 classification, output score before logistic transformation
 binary:hinge : hinge loss for binary classification. This
 makes predictions of 0 or 1, rather than producing
 probabilities.
 multi:softmax : set XGBoost to do multiclass classification
 using the softmax objective, you also need to set
 num_class(number of classes)
 multi:softprob : same as softmax, but output a vector of
 ndata nclass, which can be further reshaped to ndata nclass
 matrix. The result contains predicted probability of each
 data point belonging to each class.
 default=reg:squarederror

eval_metric – default metrics makes sense
 The metric to be used for validation data.
 The default values are rmse for regression, error for
 classification and mean average precision for ranking.

In []:

In []:

In []:

In []:

Place your answer here regarding hyperparameters.

```
In [260... import xgboost as xgb
xgb_model = xgb.XGBClassifier()
xgb_model.fit(train_X, train_Y)
y_pred = xgb_model.predict(test_X)
acc = accuracy_score(test_Y, y_pred)
print(acc)
```

0.5948504354411208

```
In [263... import xgboost as xgb
xgb_model = xgb.XGBClassifier(booster = 'dart', eta = 0.3, max_depth = 3, ob
xgb_model.fit(train_X, train_Y)
y_pred = xgb_model.predict(test_X)
acc = accuracy_score(test_Y, y_pred)
print(acc)
```

0.60734570238546

2. Now compare the XGBoost performance to the decision tree implementation from part 3. Describe in text how they compare, and if this aligns with what you expect. (3 points)

Place your code here

Decision tree was over fitting for max_depth of None, 128, 256 and 512 so I took max_depth of 16 and the test set was performing around 56.7% whereas xgboost is around 60.7%

XgBoost was bound to do better as it builds multiple trees – Each new tree is built to improve on the deficiencies of the previous trees and this concept is called boosting.

when compared to

decision tree that has only one tree so its deficiency lies there with no other chance for it to change itself

XgBoost also uses gradient of the previous tree into the new tree which helps in learning and retaining previous tree information

unlike Decision tree that has only one tree which gives less flexibility to learn

End of homework