


Recommender Systems

Aline Bessa
PhD Candidate, NYU

Modified from Jeff Ullman, Anand Rajaraman, and Jure Leskovec

<http://www.mmds.org>

I didn't know I liked that!



- **Michael**
 - Buys Metallica album
 - Buys Megadeth album



- **John**
 - Bought a Metallica album
 - Recommender system suggests Megadeth from data collected about **Michael**

pandora[®]

okcupid



Spotify[®]

Rotten
Tomatoes[®]



twitter

allrecipes[®]com

amazon

Linked in

facebook

NETFLIX

IMDb

The New York Times

hulu

yelp[®]


tripadvisor[®]

Instagram

You Tube

Relationship with Big Data

High number of users



Source: <https://www.statista.com/chart/7677/netflix-subscriber-growth/>

Relationship with Big Data

High number of items

How many products does Amazon sell (2015)?

Amazon.com	USA	488 million	+235 million
Amazon.co.uk	UK	261 million	+108 million
Amazon.de	Germany	237 million	+96 million
Amazon.fr	France	209 million	+90 million
Amazon.co.jp	Japan	168 million	+60 million
Amazon.it	Italy	165 million	+77 million
Amazon.es	Spain	160 million	+76 million
Amazon.ca	Canada	133 million	+77 million
Amazon.in	India	42 million	+18 million

Source: <https://export-x.com/2015/12/11/how-many-products-does-amazon-sell-2015/>

Relationship with Big Data

Exploration of a long tail of choices: finer tuned recommendations



Read <http://www.wired.com/wired/archive/12.10/tail.html> to learn more!

More data leads to better recommendations!

- Tastes are better characterized
- Level of personalization in recommendations increases
- More data beats better algorithms
 - <http://anand.typepad.com/datawocky/2008/03/more-data-usual.html>
- Adding more data is usually easy
 - e.g., add IMDB data on movie genres
- Cost of keeping additional data for items is cheap

Formal Model

- X = set of **users**
- S = set of **items**
- **Utility function** $u: X \times S \rightarrow R$
 - R is set of ratings
 - R is a totally ordered set
 - e.g., **0-5** stars, real number in **[0,1]**

Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Utility Matrix - Goal

	Avatar	LOTR	Matrix	Pirates
Alice	1	$u(\text{Alice}, \text{LOTR}) = ?$	0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Calculating $u(x, i)$: key problems

- Gathering “known” ratings for matrix
 - How to collect the data in the utility matrix?
 - Explicit *versus* implicit rating
- Predicting ratings $u(x, i)$ from the known ones
 - Mainly interested in **high unknown ratings**
 - We are not interested in knowing what you *don't like* but what you *like*
- Evaluating different $u(x, i)$ functions
 - How to measure success/performance of recommendation methods

Approaches to recommender systems

- Content-based
 - Recommend items to user x similar to previous items rated highly by x
- Collaborative filtering
 - Find set N of users whose ratings are similar to x 's
 - Predict new ratings for x based on ratings of users in N
 - Recommend items to x associated to highest predicted ratings
- Hybrid methods
 - Recommend items combining similarities with previous items rated by x and predicted ratings based on ratings of x 's similar users

Approaches to recommender systems

- Content-based
 - Recommend items to user x similar to previous items rated highly by x
- Collaborative filtering **TODAY!**
 - Find set N of users whose ratings are similar to x 's
 - Predict new ratings for x based on ratings of users in N
 - Recommend items to x associated to highest predicted ratings
- Hybrid methods
 - Recommend items combining similarities with previous items rated by x and predicted ratings based on ratings of x 's similar users

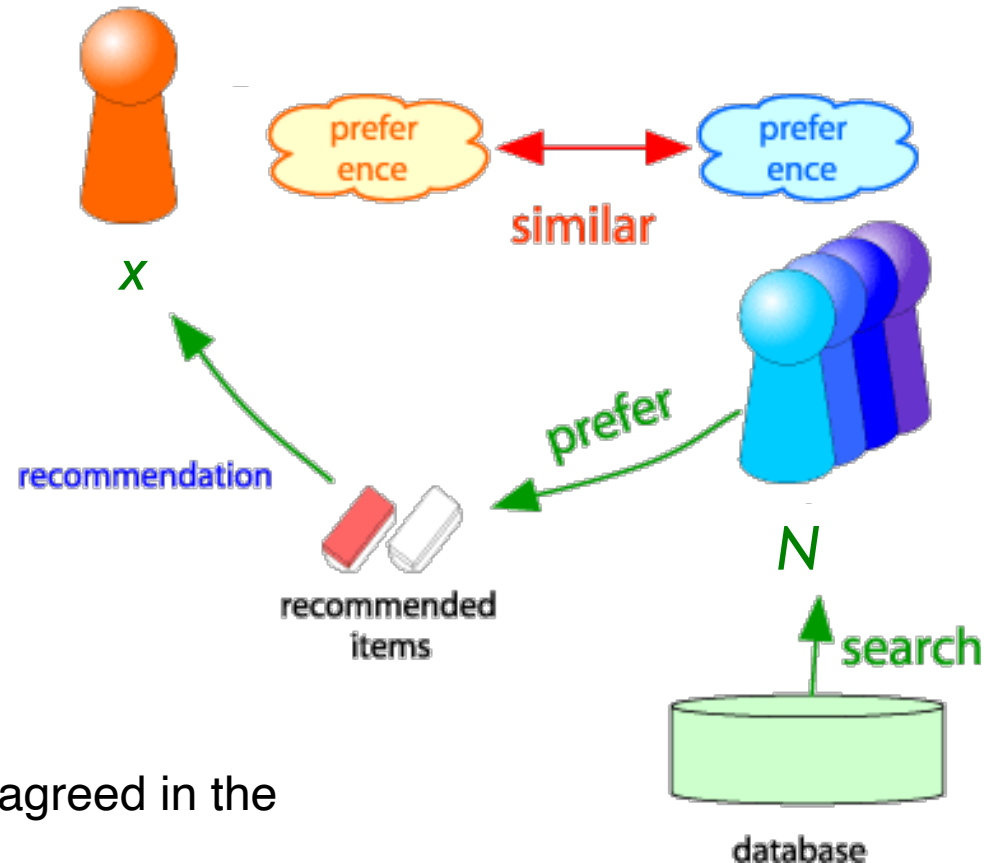
Collaborative filtering - why?

- Works for any kind of item in an agnostic fashion
 - No need to select features for similarity computations with previously rated items
 - Useful for “harder” data types (e.g., audio)
- Exploits quality judgments of other users
 - An item similar to previously rated ones can still be bad
- Significant technical advancements in the last decade
 - Netflix prize - https://en.wikipedia.org/wiki/Netflix_Prize
 - Recsys conference - <https://recsys.acm.org/best-papers/>
- An $X \times S$ rating matrix is the sole requirement
 - No additional data is used by CF algorithms

Collaborative Filtering

Consider user x

1. Find set N of other users whose ratings are “**similar**” to x ’s ratings
2. Predict x ’s ratings based on ratings of users in N



Underlying assumption I: People who agreed in the past will agree in the future

Underlying assumption II: The items people will like in the future are similar to the ones they liked in the past

1) Finding Similar Users

Input: Set of users X , value k for k most similar users

Output: Sets N of k most similar users for all x in X

For user x in X :

$N[x] = \{\}$

For user x in X :

For user y in $X \setminus \{x\}$:

tmp = compute_similarity(x , y)

if tmp is one of the k top similarities computed so far for x :

store y in $N[x]$

return N

1) Finding Similar Users

Input: Set of users X , value k for k most similar users

Output: Sets N of k most similar users for all x in X

For user x in X :

$N[x] = \{\}$

For user x in X :

For user y in $X \setminus \{x\}$:

$\text{tmp} = \text{compute_similarity}(x, y)$

if tmp is one of the k top similarities computed so far for x :

store y in $N[x]$

return N

1) Finding Similar Users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- Which users are more similar to *A*?

- Users *A* and *C*
 - Rated **two movies** in common
 - **Opposite opinions** about them
- Users *A* and *B*
 - Rated **one movie** in common
 - **Similar opinions** about it

Intuitively we want $\text{sim}(A, B) > \text{sim}(A, C)$

1) Finding Similar Users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- **Jaccard similarity**

- **Problem:** Ignores values of ratings
- $A = \{HP1, TW, SW1\}$, $B = \{HP1, HP2, HP3\}$,
 $C = \{TW, SW1, SW2\}$
- $Jaccard(A, B) = 1/5$, $Jaccard(A, C) = 1/2$
- $Jaccard(A, C) > Jaccard(A, B)$ **Intuitively wrong!**

Note:

$$Jaccard(X, Y) = |X \cap Y| / |X \cup Y|$$

1) Finding Similar Users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- **Cosine similarity**

- **Problem:** Treats missing rating as zero (low rating)
- $A = [4, 0, 0, 5, 1, 0, 0]$, $B = [5, 5, 4, 0, 0, 0, 0]$,
 $C = [0, 0, 0, 2, 4, 5, 0]$
- $\text{Cosine}(A, B) = 0.380$, $\text{Cosine}(A, C) = 0.322$
- $\text{Cosine}(A, B) > \text{Cosine}(A, C)$
- **Intuitively correct but values are close**

Note: $\text{Cosine}(X, Y) = \frac{(X \cdot Y)}{(\|X\| \cdot \|Y\|)}$

1) Finding Similar Users

- Jaccard works better when the matrix is binary
- Cosine works better when ratings are normalized
- Many other options
 - Pearson correlation
 - Kullback–Leibler divergence
 - Metric learning https://en.wikipedia.org/wiki/Similarity_learning

2) Predicting Ratings

- Let r_w be the vector of user w 's ratings
- Let N' be the subset of N where all n in N' rated item i
- The predicted rating r_{xi} (shorthand for $u(x, i)$) is

$$r_{xi} = \frac{1}{k} \sum_{y \in N'} r_{yi}$$

$$r_{xi} = \frac{\sum_{y \in N'} s_{xy} \cdot r_{yi}}{\sum_{y \in N'} s_{xy}}$$

Shorthand:

$$s_{xy} = \text{sim}(x, y)$$

What if we focus on the items instead?

- **So far:** User-user collaborative filtering
 - Similarity computed between users
- **Another view:** Item-item collaborative filtering
 - For item i , find set N of similar items
 - Predict rating for item i based on ratings for items in N
 - Can use same similarity metrics and prediction functions as in user-user model


$$r_{xi} = \frac{\sum_{j \in N'(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N'(i;x)} s_{ij}}$$

s_{ij} - similarity of items i and j
 r_{xj} - rating of user u on item j
 $N'(i;x)$ - set of items rated by x similar to i (subset of N)


Item-Item CF (INI=2)

users

	12	11	10	9	8	7	6	5	4	3	2	1	
movies		4		5			5			3		1	1
	3	1	2			4			4	5			2
		5	3	4		3		2	1		4	2	3
		2			4			5		4	2		4
	5	2					2	4	3	4			5
		4			2			3		3		1	6



- unknown rating



- rating between 1 to 5

Item-Item CF (INI=2)

		users												
		12	11	10	9	8	7	6	5	4	3	2	1	
movies			4		5			5	?		3		1	1
		3	1	2			4			4	5			2
			5	3	4		3		2	1		4	2	3
			2			4			5		4	2		4
		5	2					2	4	3	4			5
			4			2			3		3		1	6



- predict rating of movie 1 by user 5

Item-Item CF (INI=2)

		users												
		12	11	10	9	8	7	6	5	4	3	2	1	
movies			4		5			5	?		3		1	1
		3	1	2			4			4	5			2
			5	3	4		3		2	1		4	2	<u>3</u>
			2			4			5		4	2		4
		5	2					2	4	3	4			5
			4			2			3		3		1	<u>6</u>
														$\text{sim}(1,m)$
														1.00
														-0.18
														<u>0.41</u>
														-0.10
														-0.31
														<u>0.59</u>

Neighbor selection:

Identify movies similar to movie 1, rated by user 5

1) Subtract mean rating m_i from each movie i

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]

2) Compute cosine similarities between rows

Item-Item CF (INI=2)

		users													
		12	11	10	9	8	7	6	5	4	3	2	1		
movies			4		5			5	?		3		1	1	$\text{sim}(1,m)$ 1.00
		3	1	2			4			4	5			2	-0.18
			5	3	4		3		2	1		4	2	<u>3</u>	<u>0.41</u>
			2			4			5		4	2		4	-0.10
		5	2					2	4	3	4			5	-0.31
			4			2			3		3		1	<u>6</u>	<u>0.59</u>

Compute similarity weights:

$$s_{1,3}=0.41, s_{1,6}=0.59$$

Item-Item CF (INI=2)

users

	12	11	10	9	8	7	6	5	4	3	2	1	
		4		5			5	2.6		3		1	1
	3	1	2			4			4	5			2
		5	3	4		3		2	1		4	2	<u>3</u>
		2			4			5		4	2		4
	5	2					2	4	3	4			5
		4			2			3		3		1	<u>6</u>

movies

Predict by taking weighted average:

$$r_{1.5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

$$r_{ix} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$

Common Practice

- So far:

$$r_{xi} = \frac{\sum_{j \in N'(i;x)} s_{ij} r_{xj}}{\sum_{j \in N'(i;x)} s_{ij}}$$

- Common practice:

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N'(i;x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N'(i;x)} s_{ij}}$$

baseline estimate for r_{xi}

$$b_{xi} = \mu + b_x + b_i$$

μ = overall mean movie rating

b_x = rating deviation of user x

= (avg. rating of user x) - μ

b_i = rating deviation of movie i

= (avg. rating for item i) - μ

Item-Item *versus* User-User

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

- In practice, it has been observed that item-item often works better than user-user
- **Why?** Items are simpler, users have multiple tastes

Limitations of Collaborative Filtering

- Cold Start
 - New items have no ratings and new users have no history
 - CR needs enough users in the system to build N
- Sparsity
 - The user/ratings matrix is sparse
 - Hard to find users that have rated the same items
- First rater
 - Cannot recommend an item that has not been previously rated
 - Detrimental to new items or *niche* items
- Popularity bias
 - Cannot recommend items to someone with unique taste
 - Tends to recommend popular items

Evaluation

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

Evaluation

Diagram illustrating a user-movie rating matrix for evaluation.

The matrix is labeled **users** (vertical axis) and **movies** (horizontal axis).

The matrix is divided into two sections:

- Training Data (Yellow cells):** Contains known ratings.
- Test Data Set (Gray cells):** Contains unknown ratings, indicated by a blue arrow and the text "Test Data Set".

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

Evaluating Predictions

- Root mean square error (**RMSE**)
 - $\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$ where r_{xi} is the estimated rating and r_{xi}^* is the true one
- Precision at top k (**P@k**)
 - $A = \{\text{relevant items in test set for user } x\}$
 - e.g., items with ratings 4 and 5
 - $B = \{k \text{ top estimated items in test set for user } x\}$
 - $P@k = |A \cap B|/|B|$

Evaluating Predictions

- Root mean square error (**RMSE**)
 - $\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$ where r_{xi} is the estimated rating and r_{xi}^* is the true one
- Precision at top k (**P@k**)
 - $A = \{\text{relevant items in test set for user } x\}$
 - e.g., items with ratings 4 and 5
 - $B = \{k \text{ top estimated items in test set for user } x\}$
 - $P@k = |A \cap B|/|B|$

Focus on accuracy of recommendations!

Beyond Accuracy

- In the past ~10 years, there has been a growing concern about other aspects of good recommendations
 - **Coverage** - Number of items/users for which the system can make predictions
 - **Diversity** - It is more likely to find a suitable item if the recommended items are diverse, unless the user expressed narrow set of preferences
 - **Novelty** - Recommendations for items that the user did not know about are relevant
 - **Scalability** - Large collections of items in real-life scenarios require algorithmic tradeoffs (e.g., efficiency and accuracy)
 - **Co-utility** - The probability that an item has of being useful is affected by other recommended items
 - **Adaptivity** - Item collections and user interests change over time

Complexity of CF

- Finding the set **N** of most similar users/items is expensive
 - $O(|X|)$ per user or $O(|S|)$ per item

Too expensive to do at runtime!

- Pre-computation of **N** is common practice
 - Naïve pre-computation
 - Near-neighbor search in high dimensions (LSH)
 - Parallelization with MapReduce

Naïve pre-computation (item-item)

Input: Set of items I , value k for k most similar items

Output: Sets N of k most similar items for all i in I

For item i in I :

$N[i] = \{\}$

For item i in I :

For item j in $I \setminus \{i\}$:

tmp = compute_similarity(i, j)

if tmp is one of the k top similarities computed so far for i :

store j in $N[i]$

return N

Near-neighbor search (LSH)

Input: Set of items I , value k for k most similar items

Output: Sets N of k most similar items for all i in I

Note I: we abstracted the parameters for the family of hash functions in the LSH here

H = Hashtable for items

For item i in I :

$\text{compute_LSH}(H, i)$

Note II: Items a and b are in the same bucket in H , with a probability $\geq p$, if $\text{dist}(a, b) \leq r$. dist , p and r are parameters

For item i in I :

$\text{candidates} = \{\text{all items in } i\text{'s bucket in } H\} \setminus \{i\}$

$N[i] = \{\text{top } k \text{ most similar items in candidates}\}$

return N

Parallelization with MapReduce

Mapper input: Item i and a set of items J

Note: Items are either rows or columns in a utility matrix

Map function: Compute $\text{sim}(j, i)$ for all j in J ; keep **local** set of top similarities $K = \{(j_1, \text{sim}(j_1, i)), \dots, (j_k, \text{sim}(j_k, i))\}$; **output tuple** (i, K)

Reduce function: For each key i , merge the sets of its local top similarities and generate a global set K_{global} ; **output** (i, K_{global})

Ethics and recommender systems

- Acquisition of user data has to be transparent
- Privacy of users
 - <https://www.wired.com/2009/12/netflix-privacy-lawsuit/>
- Transparency in filtering choices
- Explainability of recommendations

More!

- Ricci, F., Rokach, L., Shapira, B., & Kantor, P. B. (Eds). (2011). **Recommender systems handbook**. New York: Springer.
- Mahmood, T., Ricci, F. **Learning and adaptivity in interactive recommender systems**. In: Proceedings of the 9th International Conference on Electronic Commerce, ICEC'07, pp. 75–84. ACM Press, New York, NY, USA (2007)
- Collaborative filtering and hadoop: <http://aimotion.blogspot.com/2012/08/introduction-to-recommendations-with.html>
- Winning the Netflix Prize: <http://blog.echen.me/2011/10/24/winning-the-netflix-prize-a-summary/>